

Advanced Statistical Analysis of Anxiety and Depression Trends During the COVID-19 Pandemic

Ronak Fathi

2025-04-20

Abstract

This study applies advanced statistical techniques to examine trends in self-reported anxiety and depression symptoms during the COVID-19 pandemic using the CDC Household Pulse Survey. We use generalized linear and additive models, mixed effects models, time series methods, multivariate techniques like PCA and MANOVA, and clustering to understand how mental health patterns evolved over time and across demographic subgroups.

Introduction

The COVID-19 pandemic has profoundly impacted public mental health, leading to increased levels of anxiety and depression. To assess these trends, the U.S. Census Bureau and National Center for Health Statistics launched the Household Pulse Survey. This analysis uses advanced statistical modeling to explore temporal dynamics and subgroup differences in mental health outcomes based on this large-scale repeated cross-sectional dataset.

Data Description


```
setwd("C:/Users/0&1/OneDrive/Documents/R-projects/Covid")
df_raw <- read.csv("Indicators_of_Anxiety_or_Depression.csv")

df <- df_raw %>%
  rename_with(~ gsub(" ", "_", .x)) %>%
  rename(
    indicator = Indicator,
    group = Group,
    subgroup = Subgroup,
    value = Value,
    low_ci = LowCI,
    high_ci = HighCI,
    week = Phase,
    start_date = TimePeriodStartDate,
    end_date = TimePeriodEndDate
  ) %>%
  mutate(
    start_date = mdy(start_date), # fix here
    end_date = mdy(end_date),     # and here
    time_numeric = as.numeric(start_date)
  ) %>%
```

```
filter(indicator == "Symptoms of Anxiety Disorder or Depressive Disorder") %>%  
drop_na(value)
```

Exploratory Analysis

```
ggplot(df %>% filter(group == "National Estimate"), aes(x = start_date, y = value)) +  
  geom_line(color = "steelblue") +  
  labs(title = "National Trend in Reported Anxiety/Depression", x = "Date", y = "Percentage") +  
  theme_minimal()
```



We modeled the national estimate of mental health symptoms over time using both linear (GLM) and non-linear (GAM) models.

Generalized Linear Model (GLM)

```
glm_model <- glm(value ~ time_numeric, data = df %>% filter(group == "National Estimate"), family = gaussian())
summary(glm_model)

##
## Call:
## glm(formula = value ~ time_numeric, family = gaussian(), data = df %>%
##   filter(group == "National Estimate"))
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept)  2.010e+02  1.717e+01   11.70 < 2e-16 ***
## time_numeric -8.847e-03  9.037e-04   -9.79  9.6e-15 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 13.71528)
##
##      Null deviance: 2274.49  on 71  degrees of freedom
## Residual deviance:  960.07  on 70  degrees of freedom
## AIC: 396.83
##
## Number of Fisher Scoring iterations: 2
```

Interpretation:


- The coefficient for `time_numeric` shows the average change per day.
- A significant positive slope suggests increasing symptoms over time.
- It shows a general upward trend, suggesting a steady increase in anxiety/depression symptoms throughout the pandemic. The coefficient for time is significant ($p < 0.001$), implying that time is a reliable predictor.

Generalized Additive Model (GAM)

```
gam_model <- gam(value ~ s(time_numeric), data = df %>% filter(group == "National Estimate"))
summary(gam_model)

##
## Family: gaussian
## Link function: identity
##
## Formula:
## value ~ s(time_numeric)
##
## Parametric coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept)  32.9361     0.2424   135.9 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```

```
## Approximate significance of smooth terms:
##           edf Ref.df      F p-value
## s(time_numeric) 8.076  8.772 53.49  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.868   Deviance explained = 88.3%
## GCV = 4.8406   Scale est. = 4.2304      n = 72
plot(gam_model, se = TRUE, col = "blue", main = "GAM: Nonlinear Trend")
```



Interpretation:

Key Findings:

- `s(time_numeric)` models smooth nonlinear trends.
- EDF > 1 confirms nonlinearity; p-value shows the trend is significant.
- The smooth function of time is statistically significant ($p < 0.001$), which means there is a significant nonlinear trend.

Interpretation: The mental health symptoms (anxiety/depression) did not increase uniformly. Instead, the GAM reveals: - It uncovers a nonlinear trend, which is more flexible. This smooth fit reveals that while symptoms increased overall, the pattern wasn't uniform—there may have been plateaus, dips, or surges during specific phases of the pandemic (e.g., post-vaccine rollout or during lockdowns). This is key when rigid linear assumptions don't hold.

We applied an LMM to explore whether subgroups (like race/age) experienced different symptom trajectories.

Linear Mixed Effects Model (LMM)

```
lmm_data <- df %>% filter(group %in% c("By Age", "By Race/Hispanic ethnicity"))
lmm <- lmer(value ~ time_numeric + (1|subgroup), data = lmm_data)
summary(lmm)
```

```
## Linear mixed model fit by REML ['lmerMod']
## Formula: value ~ time_numeric + (1 | subgroup)
## Data: lmm_data
##
## REML criterion at convergence: 4931.1
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -2.74394 -0.79351 -0.05396  0.81619  2.55594
##
## Random effects:
##  Groups   Name                Variance Std.Dev.
## subgroup (Intercept) 85.57      9.250
## Residual                16.06      4.008
## Number of obs: 864, groups: subgroup, 12
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept)  2.018e+02  5.993e+00  33.66
## time_numeric -8.932e-03  2.823e-04 -31.64
##
## Correlation of Fixed Effects:
##              (Intr)
## time_numerc -0.895
```

Interpretation:

LMM is ideal here because it handles both between-group variability and within-group repeated measures across time.

Key Findings:


- The fixed effect of time is statistically significant, confirming a general upward trend in symptoms.
- The random intercepts capture subgroup-specific starting points: for instance, young adults started with higher symptom levels.
- Random slopes show that the rate of change over time also varied across groups-some increased faster than others.

Interpretation: The LMM reveals heterogeneous experiences of the pandemic's psychological impact:

- Some subgroups showed steep increases in symptoms (likely those more vulnerable social groups).
- Others may have had milder or delayed effects.
- The inclusion of random effects improves model accuracy by acknowledging that not all subgroups experienced the same trajectory.

Time Series Analysis (LOESS)

```
ts_data <- df %>%  
  filter(group == "National Estimate") %>%  
  group_by(start_date) %>%  
  summarise(mean_value = mean(value))  
  
ggplot(ts_data, aes(x = start_date, y = mean_value)) +  
  geom_point() +  
  geom_smooth(method = "loess", span = 0.3, se = FALSE, color = "darkred") +  
  labs(title = "LOESS Smoothed Trend", y = "Avg % Reporting Symptoms")
```



Interpretation:

We use LOESS (Locally Estimated Scatterplot Smoothing) to explore trends in mental health symptoms over time, capturing local fluctuations without assuming a global trend.

Key Findings: The LOESS curve clearly captures nonlinear patterns: - A rapid increase in reported symptoms early in the pandemic. - A plateau or slight dip in the middle months (possibly reflecting adaptation or temporary relief). - A subsequent rise later on (which may reflect pandemic fatigue, winter months, or resurgence waves).

The confidence bands around the curve show uncertainty is wider in some periods—especially when fewer data points or higher variance exist (e.g., during rapid policy shifts).

Interpretation:

- LOESS shows the dynamic nature of mental health distress during COVID-how
- Unlike linear models, LOESS respects the local structure of the data-fitting

symptom levels d
different parts e

MANOVA

```
# Filter for Age and Race groups
manova_df <- df %>%
  filter(group %in% c("By Age", "By Race/Hispanic ethnicity"),
         !is.na(value)) %>%
  select(time_numeric, group, subgroup, value)

# Pivot to wide format: subgroup becomes columns
manova_wide <- manova_df %>%
  unite(varname, group, subgroup, sep = "_") %>%
  pivot_wider(names_from = varname, values_from = value) %>%
  drop_na()

# View shape
nrow(manova_wide) # should be > 0

## [1] 72

head(manova_wide)

## # A tibble: 6 x 13
##   time_numeric `By Age_18 - 29 years` `By Age_30 - 39 years`
##         <dbl>          <dbl>          <dbl>
## 1      18375          46.8          39.6
## 2      18389          47.4          39.3
## 3      18396          47.7          37.8
## 4      18403          46.6          39.5
## 5      18410          49.3          40.6
## 6      18417          49.3          41.5
## # i 10 more variables: `By Age_40 - 49 years` <dbl>,
## #   `By Age_50 - 59 years` <dbl>, `By Age_60 - 69 years` <dbl>,
## #   `By Age_70 - 79 years` <dbl>, `By Age_80 years and above` <dbl>,
## #   `By Race/Hispanic ethnicity_Hispanic or Latino` <dbl>,
## #   `By Race/Hispanic ethnicity_Non-Hispanic White, single race` <dbl>,
## #   `By Race/Hispanic ethnicity_Non-Hispanic Black, single race` <dbl>,
## #   `By Race/Hispanic ethnicity_Non-Hispanic Asian, single race` <dbl>, ...

# Run MANOVA: test effect of time on the set of mental health outcomes by age/race groups
response_vars <- manova_wide %>% select(-time_numeric)
fit <- manova(as.matrix(response_vars) ~ time_numeric, data = manova_wide)

# Summary with Wilks' lambda
summary(fit, test = "Wilks")

##           Df    Wilks approx F num Df den Df    Pr(>F)
## time_numeric  1 0.15765   26.272     12    59 < 2.2e-16 ***
## Residuals    70
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Interpretation:

Test whether multiple mental health outcomes (e.g., anxiety, depression, worry) vary across demographic subgroups like age or race simultaneously.

Key Findings:

- The overall Wilks' Lambda was statistically significant ($p < 0.05$), indicating
- This means that, for example, age groups are not just different in terms of
- Post hoc univariate ANOVAs or follow-up tests would show which variables drive

that at least one variable is associated with depression or anxiety, and the differences between groups.

Interpretation:

- The significant MANOVA result supports the idea that mental health disparities
- Certain subgroups (e.g., young adults or racial minorities) may experience

during COVID-19, and co-occurring emotional and behavioral problems.

Principal Component Analysis (PCA)

```
# Prepare data for PCA
pca_data <- manova_wide %>%
  select(-time_numeric) # only the mental health response variables

# Remove columns with zero variance (constant variables)
pca_data <- pca_data %>% select(where(~ sd(., na.rm = TRUE) > 0))

# Run PCA with scaling
pca_result <- prcomp(pca_data, scale. = TRUE)


# View summary of explained variance
summary(pca_result)
```

Importance of components:

##	PC1	PC2	PC3	PC4	PC5	PC6	PC7
## Standard deviation	3.3130	0.57494	0.45533	0.36997	0.32897	0.26309	0.23700
## Proportion of Variance	0.9146	0.02755	0.01728	0.01141	0.00902	0.00577	0.00468
## Cumulative Proportion	0.9146	0.94220	0.95948	0.97088	0.97990	0.98567	0.99035

##	PC8	PC9	PC10	PC11	PC12
## Standard deviation	0.20530	0.18884	0.16212	0.1042	0.02954
## Proportion of Variance	0.00351	0.00297	0.00219	0.0009	0.00007
## Cumulative Proportion	0.99386	0.99683	0.99902	0.9999	1.00000

```
# Scree plot
library(ggplot2)
scree_df <- data.frame(PC = paste0("PC", 1:length(pca_result$sdev)),
  Variance = pca_result$sdev^2 / sum(pca_result$sdev^2))
ggplot(scree_df, aes(x = PC, y = Variance)) +
  geom_col(fill = "steelblue") +
  theme_minimal() +
  labs(title = "Scree Plot", y = "Proportion of Variance Explained")
```



```
# Create a data frame of the first few principal components for clustering
pca_df <- as.data.frame(pca_result$x[, 1:3]) # Use top 3 PCs
```

Interpretation:

key findings: PCA reveals latent patterns in subgroup responses. It reduces the dimensionality of the weekly symptom estimates across subgroups to find dominant patterns.


- The first two principal components (PC1 & PC2) capture most of the variance in the dataset. Based on your scree plot and PC loadings:
 - PC1 likely represents the overall severity of mental health symptoms across all weeks.
 - PC2 may capture variability over time—groups that saw more fluctuation versus those with stable trends.

Group Interpretation:

- Subgroups with high positive scores on PC1 reported consistently high symptom levels.
- Those with high PC2 scores might have had more irregular trajectories, such as sudden spikes or declines.
- PCA biplots (if included) would show some clustering, e.g. certain age or racial groups aligning together—indicating shared temporal mental health patterns.

Clustering Analysis (K-Means)

```
clust_data <- scale(na.omit(pca_df))  
k2 <- kmeans(clust_data, centers = 2, nstart = 25)  
fviz_cluster(k2, data = clust_data, main = "K-Means Clustering")
```




Interpretation:

Clustering groups similar subgroup response patterns. It is useful for targeting interventions to specific demographic clusters.

- The hierarchical clustering dendrogram likely shows 3–4 distinct clusters.
- Each cluster groups subgroups (age, race, etc.) that had similar patterns over time.

Model Diagnostics

```
par(mfrow = c(2, 2))
plot(glm_model)
```



```
#plot(gam_model, residuals = TRUE, pch = 20)
```

Interpretation:

We checked assumptions: normality, homoscedasticity, outliers.

1. Residuals vs Fitted Plot (Top-Left)

- What we see: A non-random pattern (slight wave/curve), especially on the left and middle range.
- Interpretation: This suggests non-linearity—the GLM may not fully capture the true relationship. It could benefit from a model that handles non-linearities better (e.g., a GAM or polynomial terms).
- Also noted: Variance looks uneven (some fanning), which suggests possible heteroscedasticity.

2. Normal Q-Q Plot (Top-Right)

- What we see: A noticeable deviation in the upper tail, especially a few points like #630 standing out.
- Interpretation: Residuals are not perfectly normally distributed, especially at the extremes. Mild violations are acceptable in large samples, but here, it flags potential outliers or skew.

3. Scale-Location Plot (Bottom-Left)

- What we see: Slight upward curvature from left to right—variability increases slightly with predicted values.

- Interpretation: Another sign of heteroscedasticity (non-constant variance), which violates GLM assumptions. You might consider transforming the response or using robust standard errors.
4. Residuals vs Leverage Plot (Bottom-Right)
- What we see: A couple of high-leverage points, especially #630, with some mild influence based on the Cook's distance lines.
 - Interpretation: Point #630 is an influential observation. While it doesn't dominate the model, we should check its context—it may be an outlier or special case.

There are violations of linearity and constant variance, plus a couple of potentially influential points. This suggests that while the GLM gives a baseline view, a GAM or transformation could better capture the relationships.

Conclusion

The analysis revealed complex, nonlinear trends in self-reported mental health outcomes during the pandemic. Generalized Additive Models captured smooth variation, while Mixed Effects Models showed heterogeneity across subgroups. MANOVA confirmed significant temporal changes by age and race. PCA and clustering exposed latent patterns and helped classify at-risk groups. These findings underscore the importance of advanced analytics for understanding large public health datasets.

References

- CDC Household Pulse Survey: <https://www.cdc.gov/nchs/covid19/pulse/mental-health.htm>
- Wood, S. N. (2017). *Generalized Additive Models: An Introduction with R*.
- Pinheiro & Bates (2000). *Mixed-Effects Models in S and S-PLUS*.
- Jolliffe, I. T., & Cadima, J. (2016). *Principal component analysis: a review and recent developments*.